

Project 3: Real-time 2-D Object Recognition

- Due Wednesday by 2:59am
- Points 30

This project is about 2D object recognition. The goal is to have the computer identify a specified set of objects placed on a white surface in a translation, scale, and rotation invariant manner from a camera looking straight down. The computer should be able to recognize single objects placed in the image and show the position and category of the object in an output image. If provided a video sequence, it should be able to do this in real time.

Setup

For development, you can use this [development image set](https://northeastern.instructure.com/courses/192426/files/28975733?wrap=1)

(<https://northeastern.instructure.com/courses/192426/files/28975733?wrap=1>) 

(https://northeastern.instructure.com/courses/192426/files/28975733/download?download_frd=1) . Once you have your system working, you'll need to create your own workspace for real-time recognition. If you have a webcam, the easiest thing to do is set it up on some kind of stand facing down onto a flat surface covered with white paper. If you have only a built-in laptop camera, find a section of white wall in front of which you can hold the objects to be identified. Alternatively, you may be able to stream your phone's camera video to your laptop. If you can get a setup where with a downward facing camera on a white surface, that is the easiest thing to do. If none of those work for you, capture some videos or still images with a phone and set up your program to process single images or video sequences.

Your system needs to be able to differentiate at least 5 objects of your choice. Pick objects that are differentiable by their 2D shape and are mostly a uniform dark color, similar to the dev image set. The goal of this project is a real-time system, but if that is not feasible given your setup, you can have your system take in a directory of images or a pre-recorded video sequence and process them instead. Make sure it is clear in your report which type of system you build.

If you are working on your own, you must write the code from scratch for at least one of the first four tasks--thresholding, morphological filtering, connected components analysis, or moment calculations. Otherwise, you may use OpenCV functions. If you are working in pairs, you must write the code for two of the first four tasks from scratch.

Tasks

1. **Threshold the input video**

Using the video framework from the first project (if you wish), start building your OR system by implementing a thresholding algorithm of some type that separates an object from the background. Give your system the ability to display the thresholded video (remember, you can create multiple output windows using OpenCV). Test it on the complete set of objects to be recognized to make sure this step is working well. In general, the objects you use should be darker than the background, and the whole area should be well-lit.

You may want to pre-process the image before thresholding. For example, blurring the image a little can make the regions more uniform. You could also consider making strongly colored pixels (pixels with a high saturation value) be darker, moving them further away from the white background (which is unsaturated).

It may also be helpful to dynamically set your threshold by analyzing the pixel values. Running a k-means algorithm on a random sample of pixel values (e.g. using 1/16 of the pixels in the image) with $K=2$ (also called the ISODATA algorithm) gives you the means of the two dominant colors in the image. The value in between those means may be a good threshold to use.

Required images: Include 2-3 examples of the thresholded objects in your report

2. Clean up the binary image

Clean up your thresholded image using some type of morphological filtering. Look at your images and decide whether they are displaying noise, holes, or other issues. Pick a strategy using morphological filtering (for example, growing/shrinking) to try to solve the issue(s). Explain your reasoning in your report.

Required images: Include 2-3 examples of cleaned up images in your report. Use the same images as displayed for task 1.

3. Segment the image into regions

The next step is to run connected components analysis on the thresholded and cleaned image to get regions. Give your system the ability to display the regions it finds. A good extension is to enable recognition of multiple objects simultaneously (but this is not required). Your system should ignore any regions that are too small, and it can limit the recognition to the largest N regions.

A segmentation algorithm will return a region map. If you try to display the region map directly, it won't tell you much because the region IDs will all be similar. A better method of displaying the region map is to ignore regions that are too small and to pick from a color palette for the remaining regions (e.g. a list of 256 random colors). Once you have removed regions that are too small, it can be helpful to renumber the remaining regions in sequential order. If you want to avoid color flickering in your region map display, you will have to remember the centroid and color for each region in the prior image and try to match centroids and colors in the next image.

OpenCV has a connected components function. If you don't write your own, use the one that also returns the stats. If you chose to write your own, both the two-pass connected components algorithm and the region growing methods are good algorithms to know.

A good strategy to use when getting started is to focus on the region that is larger than a certain size, is most central to the image, and does not touch the image boundary. The region size, region centroid, and region axis-oriented bounding box are all you need to make this comparison. Those statistics are easy to compute from a region map, and they are provided by the built-in OpenCV `connectedComponentsWithStats` function.

Required images: Include 2-3 examples of region maps, ideally with each region shown in a different color. Use the same images as for the prior tasks.

4. Compute features for each major region

Write a function that computes a set of features for a specified region given a region map and a region ID. Start by calculating the axis of least central moment and the oriented bounding box. Display these overlaying the objects in your output.

Important: for this assignment, the purpose is to analyze regions, not boundaries. Use region-based analysis, not boundary analysis, for generating features.

Percent filled and the bounding box height/width ratio are good features to implement first. Any features you use should be translation, scale, and rotation invariant such as moments around the central axis of rotation. Give your system the ability to display at least one feature in real time on the video output. Then you can easily test whether the feature is translation, scale, and rotation invariant by moving around the object and watching the feature value. Start with just 2-3 features and add more later. There is an OpenCV function that will compute a set of moments. If you use it, explain what moments it is calculating and which ones you are using as features in your report. Whatever features you choose to use, explain what they are in your report and show the feature vector for the images included in the report for this task.

Required images: Include 2-3 examples of regions showing the axis of least central moment and the oriented bounding box. use the same images as for the prior tasks. Include the computed feature vectors for each object.

5. Collect training data

Enable your system to collect feature vectors from objects, attach labels, and store them in an object DB (e.g. a file). In other words, your system needs to have a training mode that enables you to collect the feature vectors of known objects and store them for later use in classifying unknown objects. You may want to implement this as a response to a key press: when the user types an N, for example, have the system prompt the user for a name/label and then store the feature vector for the current object along with its label into a file. This could also be done from labeled still images of the object, such as those from a training set.

Explain in your report how your training systems works.

6. Classify new images

Enable your system to classify a new feature vector using the known objects database and a scaled Euclidean distance metric $[(x_1 - x_2) / \text{stdev}_x]$. Feel free to experiment with other distance metrics. Label the unknown object according to the closest matching feature vector in the object DB (nearest-neighbor recognition). Have your system indicate the label of the object on the output video stream. An extension is to detect when an unknown object (something not in the object DB) is in the video stream or provided as a single image.

Required images: include a result image for each category of object in your report with the assigned label clearly shown.

7. Evaluate the performance of your system

Evaluate the performance of your system on at least 3 different images of each object in different positions and orientations. Build a 5x5 confusion matrix of the results showing true labels versus classified labels. *Include the confusion matrix in your report.* Hint: you can add hooks into your code to make this easier if you can tell the system if it is correct.

8. Capture a demo of your system working

Take a video of your system running and classifying objects. If your system works on still images, show the visualizations your system makes. Include a link to the video in your report.

9. Implement a second classification method

Choose from the following options.

- A. Implement a different classifier--something besides nearest neighbor--of your choice. For example, implement K-Nearest Neighbor matching with $K > 1$. Note, KNN matching requires multiple training examples for each object and should **not** use the voting method.
- B. Use NN matching, but compare different distance metrics. For example, compare simple Euclidean, scaled Euclidean, cosine distance, scaled L-1, or scaled L-inf.
- C. Build a decision tree, either by hand, or try something like the scikit-learn python package to build a 2 or 3 level tree and then implement it as a set of if statements in your C++ code.
- D. Using the oriented bounding box, convert the pixels in the box into a fixed-size square and use the resulting square binary image as the feature vector using SSD as the distance metric.

Explain your choice and how you implemented it. Also, compare the performance with your baseline system and include the comparison in your report.

Extensions


- Write a better GUI to show your system in action and to manage the items to your DB
 - Write more than one (or two if working in pairs) of the stages of the system from scratch.
 - Add more than the required five objects to the DB so your system can recognize more objects.
 - Enable your system to learn new objects automatically by first detecting whether an object is known or unknown, and then collecting statistics for it if the object is unknown. Demonstrate how you can quickly develop a DB using this feature.
 - Experiment with more classifiers and/or distance metrics in the final task.
 - Explore some of the object recognition tools in OpenCV.
-

Report

When you are done with your project, write a short report that demonstrates the functionality of each task. You can write your report in the application of your choice, but you need to submit it **as a pdf** along with your code. Your report should have the following structure. Please **do not include code** in your report.

1. A short description of the overall project in your own words. (200 words or less)
 2. Any required images along with a short description of the meaning of the image.
 3. A description and example images of any extensions.
 4. A short reflection of what you learned.
 5. Acknowledgement of any materials or people you consulted for the assignment.
-

Submission

Submit your code and report to **Gradescope**  (<https://www.gradescope.com>). When you are ready to submit, upload your code, report, and a readme file. The readme file should contain the following information.

- Your name and any other group members, if any.
- Links/URLs to any videos you created and want to submit as part of your report.
- What operating system and IDE you used to run and compile your code.
- Instructions for running your executables.
- Instructions for testing any extensions you completed.
- Whether you are using any time travel days and how many.

For project 3, submit your .cpp and .h (.hpp) files, pdf report, and readme.txt (readme.md). Note, if you find any errors or need to update your code, you can resubmit as many times as you wish up until the deadline.

As noted in the syllabus, projects submitted by the deadline can receive full credit for the base project and extensions. (max 30/30). Projects submitted up to a week after the deadline can receive full credit for the base project, but not extensions (max 26/30). You also have eight time travel days you can use during the semester to adjust any deadline, using up to three days on any one assignment (no fractional days). If you want to use your time travel days, indicate that in your readme file. If you need to make use of the "stuff happens" clause of the syllabus, contact the instructor as soon as possible to make alternative arrangements.

Receiving grades and feedback

After your project has been graded, you can find your grade and feedback on Gradescope. Pay attention to the feedback, because it will probably help you do better on your next assignment.